

Zircon Kernel Introduction

2018/01/16

Brief Introduction

- The kernel and core platform of Google Fuchsia OS
- Derived from LK (Little Kernel) but quite different nowadays
- 64bit only system (support X86-64 and ARM64)
- Microkernel with capability-based security model
- Target modern phones and modern PC, not small system with very limited hardware resources (LK, FreeRTOS and ThreadX is suitable for small system)
- Reference:
 - <https://fuchsia.googlesource.com/zircon/+/HEAD/docs/>

Kernel Objects and Syscalls

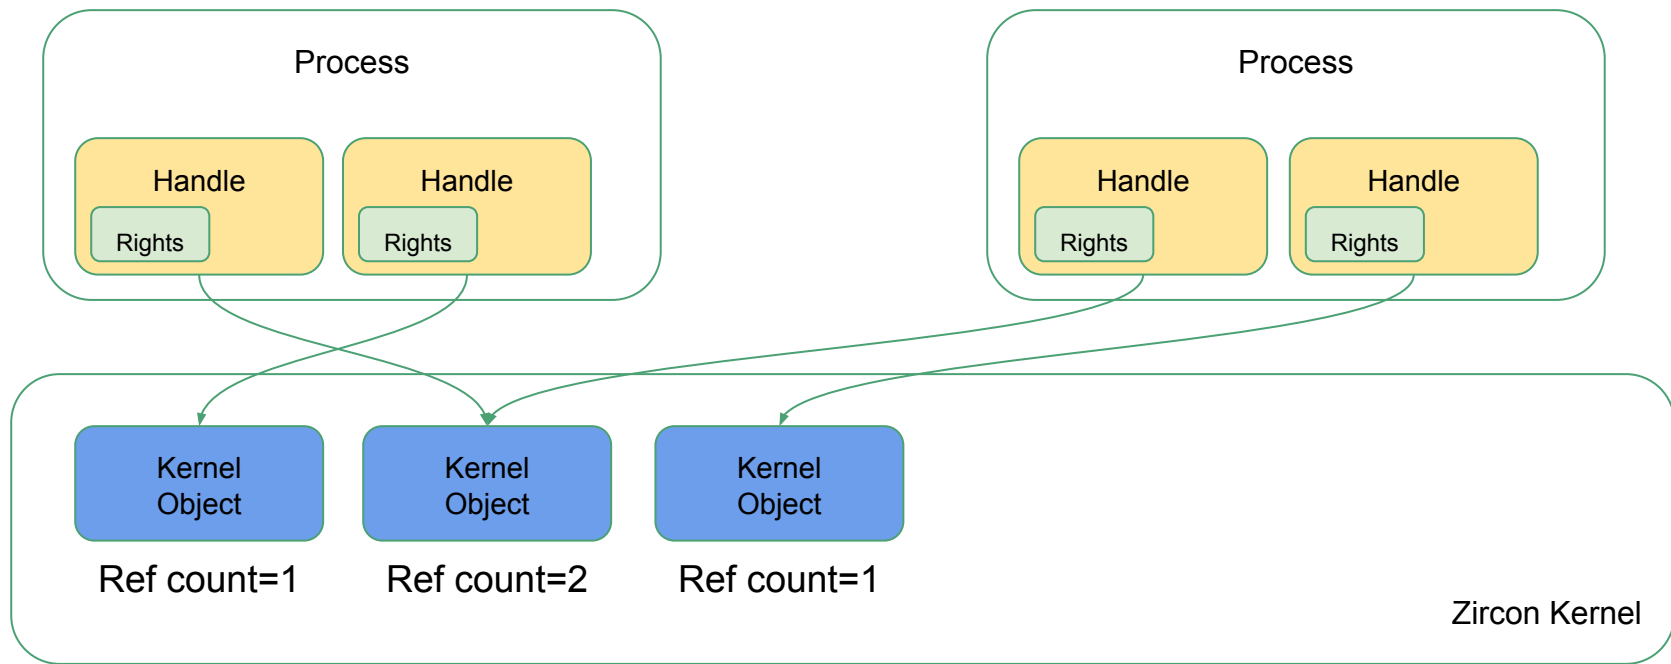
Kernel Objects

- Zircon kernel manages different types of kernel objects
 - **IPC:** Channel, Socket, FIFO
 - **Tasks:** Process, Thread, Job, Task
 - **Signaling:** Event, Event Pair, Futex
 - **Memory:** VMO, VMAR, BTI
 - **Wait:** Port
 - **For drivers:** Interrupt, Resource
- Kernel objects are C++ classes which implement the “Dispatcher” interface
- Userspace code interacts with kernel objects via **system calls** and **handles**
- Dispatcher types
 - **SoloDispatcher:** single endpoint, e.g. Event, Process, Thread, VMO
 - **PeeredDispatcher:** has two endpoints, e.g. Channel, Socket, FIFO, Event Pair

Handles

- Can be thought as a session or connection to a particular kernel object
- Can be bound to **process** or **kernel** (in-transit)
- Each process has its own handle table maintained by kernel
- A kernel object may have multiple handles that refer to it
- When the last open handle that refer to an kernel object is closed
 - The object is destroyed or,
 - The kernel marks the object for garbage collection
- Handle may be transferred to another process by
 - Writing handle into a **Channel** via `zx_channel_write()`
 - Passing handle as the 'arg1' parameter in `zx_process_start()` to the first thread in a new process
- Handle rights
 - Specify what operations on the kernel object are allowed
 - Different handles to the same kernel object can have different rights

Object, Handle and Rights

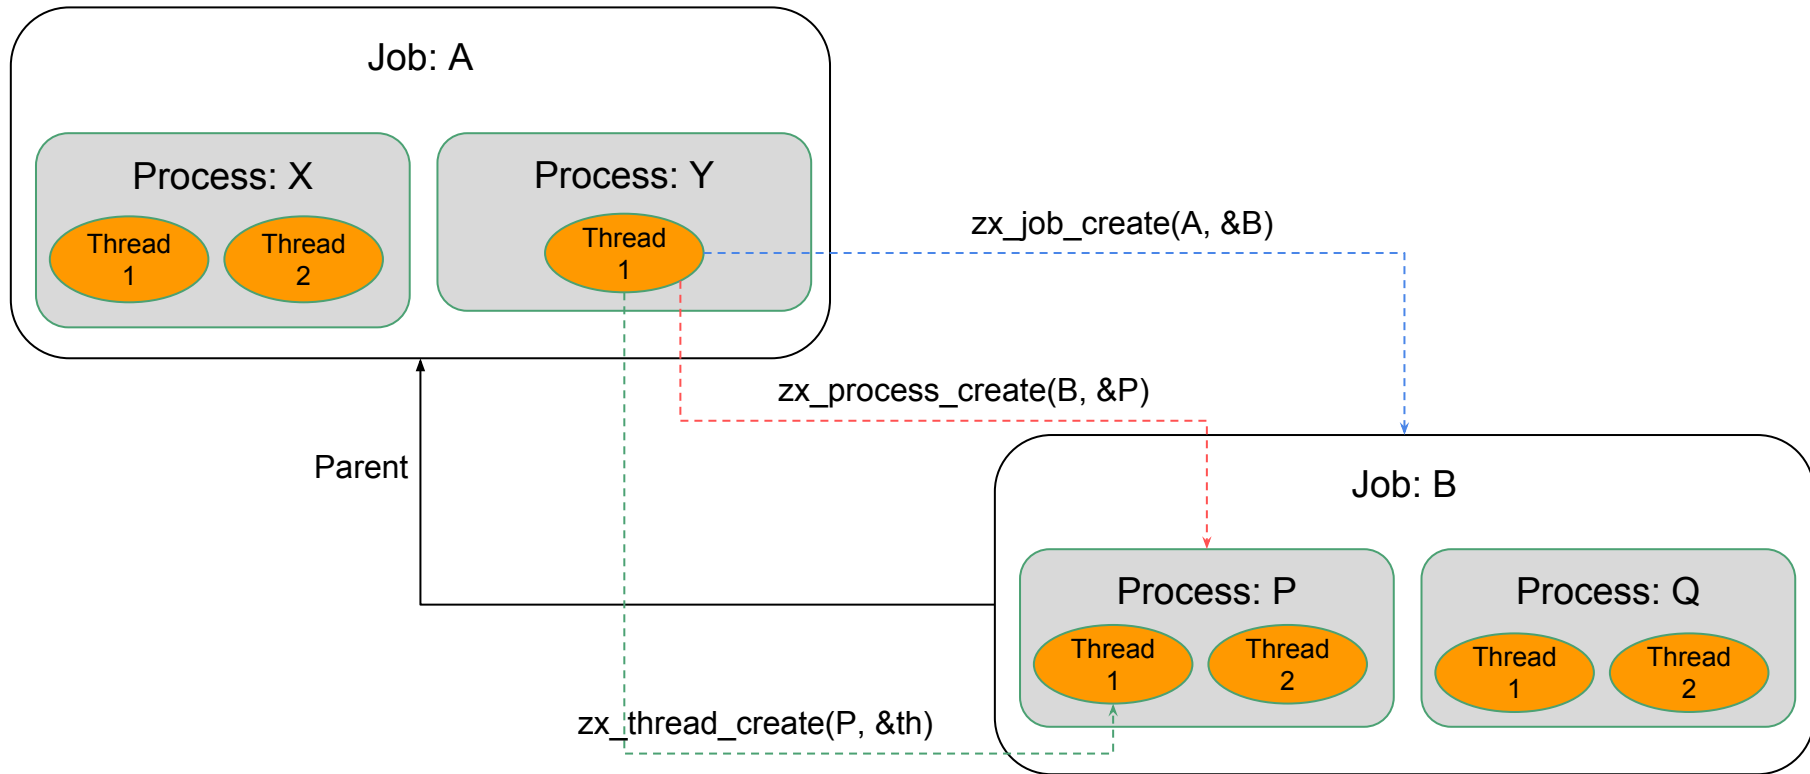


System Calls


- For userspace code to interact with kernel objects
- System calls fall into three broad categories from an access standpoint:

Access Type	Permission Check	Example
No limitation	No	<code>zx_clock_get()</code> <code>zx_nanosleep()</code>
Interact with specific object (a handle as first parameter)	<ul style="list-style-type: none">• handle is valid• handle is of correct type• handle has required rights	<code>zx_channel_write()</code> <code>zx_port_queue()</code>
Create new object (no handle provided)	Job policy	<code>zx_channel_create()</code> <code>zx_event_create()</code>

Jobs, Processes, and Threads



IPC: Channel, Socket, and FIFO



- Creating an IPC object will return two handles, one referring to each endpoint of the object

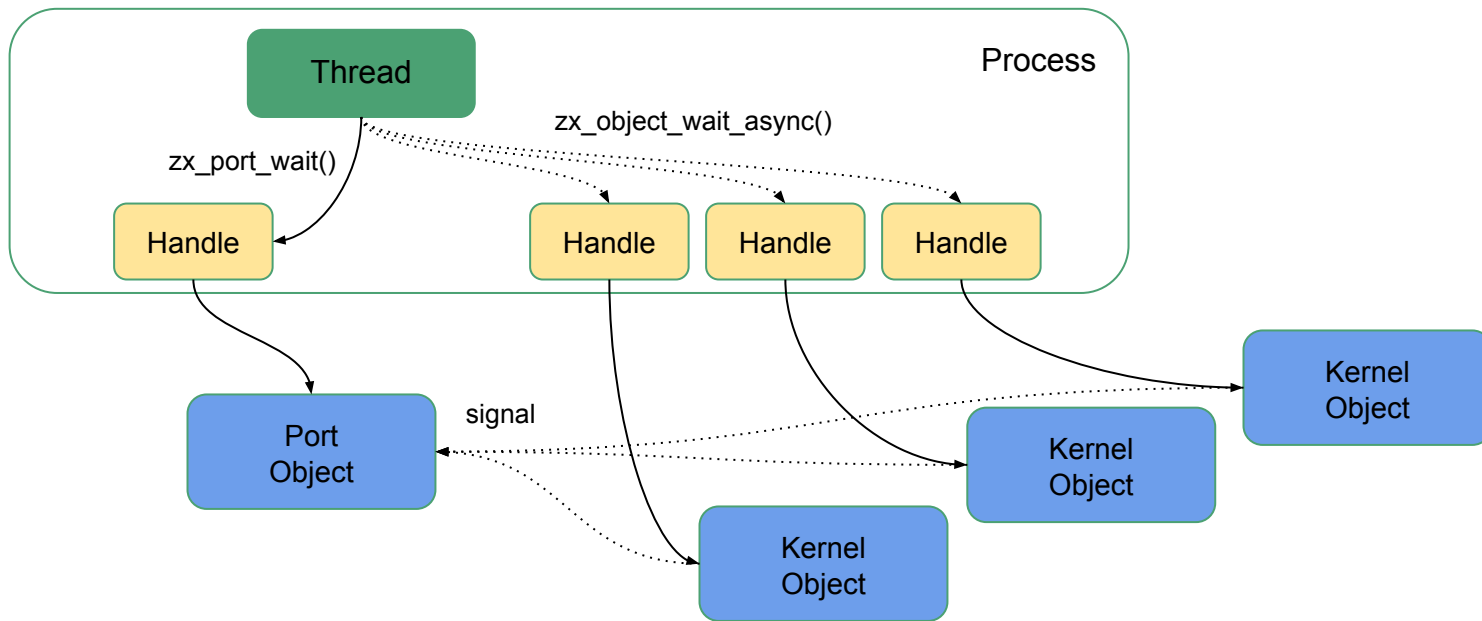
	Transfer Behavior Comparison	Handle Transfer?	Efficiency	Limitation
Socket	<ul style="list-style-type: none">• Stream-oriented• Short read/write are allowed• Data Units: bytes	No	Low	
Channel	<ul style="list-style-type: none">• Datagram-oriented• Short read/write are NOT allowed• Data Units: bytes	Yes	Mid	Single Write: Max Data Bytes: 64KB Max Handles: 64
FIFO	<ul style="list-style-type: none">• Datagram-oriented• Data Units: element size (assigned when <code>zx_fifo_create()</code>)	No	High	Single Write: Max Data Bytes: 4KB

Object Signals and Waiting

- Objects may have up to 32 signals expose to application to represent their current state
 - e.g. `ZX_CHANNEL_READABLE`, `ZX_PROCESS_TERMINATED`
- Threads may wait for signals to become active on one or more objects
- Syscalls for waiting signals: (handle should have `ZX_RIGHT_WAIT`)
 - **`zx_object_wait_one()`**: Blocking wait on a object
 - **`zx_object_wait_many()`**: Blocking wait on one or more objects (up to 8)
 - **`zx_object_wait_async()`** + **`zx_port_wait()`**: Blocking wait on one or more objects (> 8)
- User signals
 - **`ZX_USER_SIGNAL_0`** through **`ZX_USER_SIGNAL_7`**
 - User signals can be asserted or deasserted by **`zx_object_signal()`** or **`zx_object_signal_peer()`**
 - Handle should have `ZX_RIGHT_SIGNAL` or `ZX_RIGHT_SIGNAL_PEER`

Port

- If a thread is going to wait on a **large set** of handles, it is more efficient to use a Port
- Multiple signal can be accumulated at port and retrieve at once
- Signal can be queued in port, even no thread is waiting on it



Shared Memory: Virtual Memory Object (VMO)

- Used to represent both paged and physical memory
- The standard method of sharing memory between
 - Processes
 - The kernel and userspace
- **zx_vmo_read(), zx_vmo_write():** for userspace to perform basic I/O directly
- **zx_vmo_set_size():** adjust the size of VMO
- **zx_vmar_map():** map VMO in to process address space
- **zx_vmo_op_range():** for some low level operations
 - commit, decommit range of pages
 - cache sync, clean and invalidate
- **zx_vmo_set_cache_policy():** change VMO cache policy

Address Space Management: Virtual Memory Address Region (VMAR)

- Used by the kernel and userspace to represent the allocation of an address space
- Each process starts with a **root VMAR** that spans the entire address space
- Can be logically divided into any number of **non-overlapping** child VMARs
- Have a hierarchical permission model (child \leq parent)
- All allocations of address space are randomized by default
- Syscalls
 - **zx_vmar_allocate()**, **zx_vmar_destroy()**: create or destroy a VMAR and its child VMARs
 - **zx_vmar_map()**, **zx_vmar_unmap()**: map or unmap a VMO into or from a VMAR
 - **zx_vmar_protect()**: adjust memory access permissions

Fast Userspace Mutex (Futex)

- A low level synchronization primitive as a building block for higher level APIs
 - e.g. pthread APIs for mutexes, condition variables
- Zircon futexes are strictly a process local concept and cannot be shared across address spaces
- Syscalls
 - **zx_futex_wait()**: wait on a futex
 - **zx_futex_wake()**: wake some number of threads waiting on a futex
 - **zx_futex_requeue()**:
 - wake some number of threads waiting on a futex
 - move more waiters to another wait queue

Resource

- A type of kernel object for controlling access to:
 - Specific range of address space (ZX_RSRC_KIND_MMIO)
 - IRQ numbers (ZX_RSRC_KIND_IRQ)
 - Hypervisor (ZX_RSRC_KIND_HYPERVISOR)
 - SMC (ZX_RSRC_KIND_SMC)
 - ...
- Resource objects are typically private to the DDK and platform bus drivers
- ZX_RSRC_KIND_ROOT resource object is required for creating other kind of resource objects
- Resource allocations can be either shared or exclusive (owned by the holder)

Security Control: Job Policy, Handle Rights and Resources

```
// Input structure to use with ZX_JOB_POL_BASIC.
typedef struct zx_policy_basic {
    uint32_t condition;
    uint32_t policy;
} zx_policy_basic_t;

// Conditions handled by job policy.
#define ZX_POL_BAD_HANDLE 0u
#define ZX_POL_WRONG_OBJECT 1u
#define ZX_POL_VMAR_WX 2u
#define ZX_POL_NEW_ANY 3u
#define ZX_POL_NEW_VMO 4u
#define ZX_POL_NEW_CHANNEL 5u
#define ZX_POL_NEW_EVENT 6u
#define ZX_POL_NEW_EVENTPAIR 7u
#define ZX_POL_NEW_PORT 8u
#define ZX_POL_NEW_SOCKET 9u
#define ZX_POL_NEW_FIFO 10u
#define ZX_POL_NEW_TIMER 11u
#define ZX_POL_NEW_PROCESS 12u
#ifdef KERNEL
#define ZX_POL_MAX 13u
#endif

// Policy actions.
// ZX_POL_ACTION_ALLOW and ZX_POL_ACTION_DENY can
// ZX_POL_ACTION_KILL implies ZX_POL_ACTION_DENY.
#define ZX_POL_ACTION_ALLOW 0u
#define ZX_POL_ACTION_DENY 1u
#define ZX_POL_ACTION_EXCEPTION 2u
#define ZX_POL_ACTION_KILL 5u
```

```
#define ZX_RIGHT_NONE ((zx_rights_t)0u)
#define ZX_RIGHT_DUPLICATE ((zx_rights_t)1u << 0)
#define ZX_RIGHT_TRANSFER ((zx_rights_t)1u << 1)
#define ZX_RIGHT_READ ((zx_rights_t)1u << 2)
#define ZX_RIGHT_WRITE ((zx_rights_t)1u << 3)
#define ZX_RIGHT_EXECUTE ((zx_rights_t)1u << 4)
#define ZX_RIGHT_MAP ((zx_rights_t)1u << 5)
#define ZX_RIGHT_GET_PROPERTY ((zx_rights_t)1u << 6)
#define ZX_RIGHT_SET_PROPERTY ((zx_rights_t)1u << 7)
#define ZX_RIGHT_ENUMERATE ((zx_rights_t)1u << 8)
#define ZX_RIGHT_DESTROY ((zx_rights_t)1u << 9)
#define ZX_RIGHT_SET_POLICY ((zx_rights_t)1u << 10)
#define ZX_RIGHT_GET_POLICY ((zx_rights_t)1u << 11)
#define ZX_RIGHT_SIGNAL ((zx_rights_t)1u << 12)
#define ZX_RIGHT_SIGNAL_PEER ((zx_rights_t)1u << 13)
#define ZX_RIGHT_WAIT ((zx_rights_t)1u << 14)
#define ZX_RIGHT_INSPECT ((zx_rights_t)1u << 15)
#define ZX_RIGHT_MANAGE_JOB ((zx_rights_t)1u << 16)
#define ZX_RIGHT_MANAGE_PROCESS ((zx_rights_t)1u << 17)
#define ZX_RIGHT_MANAGE_THREAD ((zx_rights_t)1u << 18)
#define ZX_RIGHT_APPLY_PROFILE ((zx_rights_t)1u << 19)
#define ZX_RIGHT_SAME_RIGHTS ((zx_rights_t)1u << 31)
```

```
#define ZX_RSRC_KIND_MMIO ((zx_rsrc_kind_t)0u)
#define ZX_RSRC_KIND_IRQ ((zx_rsrc_kind_t)1u)
#define ZX_RSRC_KIND_IOPORT ((zx_rsrc_kind_t)2u)
#define ZX_RSRC_KIND_HYPERVISOR ((zx_rsrc_kind_t)3u)
#define ZX_RSRC_KIND_ROOT ((zx_rsrc_kind_t)4u)
#define ZX_RSRC_KIND_VMEX ((zx_rsrc_kind_t)5u)
#define ZX_RSRC_KIND_SMC ((zx_rsrc_kind_t)6u)
#define ZX_RSRC_KIND_COUNT ((zx_rsrc_kind_t)7u)
```


Virtual Dynamic Shared Object (vDSO)

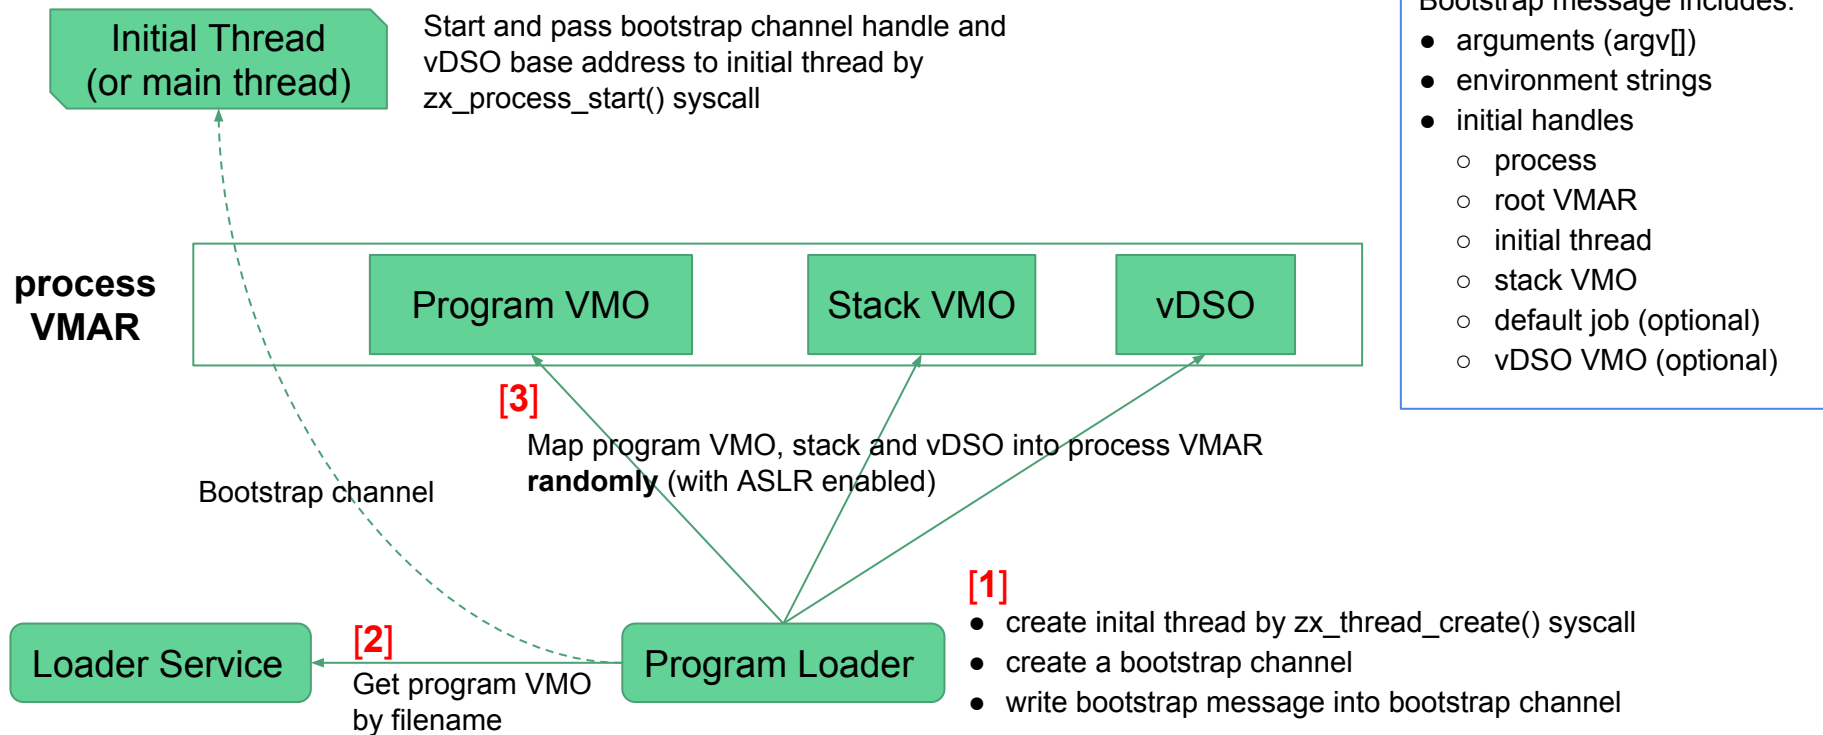
- vDSO is the only way to access to system calls for userspace
- vDSO image (ELF format) is provided directly by the kernel
- The kernel exposes vDSO to userspace as a read-only VMO
- When a new process created by program loader, the program loader:
 - Maps the vDSO into the new process's address space
 - Passes the vDSO base address to the first thread of the new process
 - Passes a vDSO VMO handle to the new process (PA_VMO_VDSO)
- The kernel enforces correct use of the vDSO in two ways:
 - It constrains how the vDSO VMO can be mapped into a process
 - vDSO VMO can only be mapped once in process address space
 - Once the vDSO mapping has been established in a process, it cannot be removed
 - It constrains what PC locations can be used to enter the kernel

Program Loading and Userboot

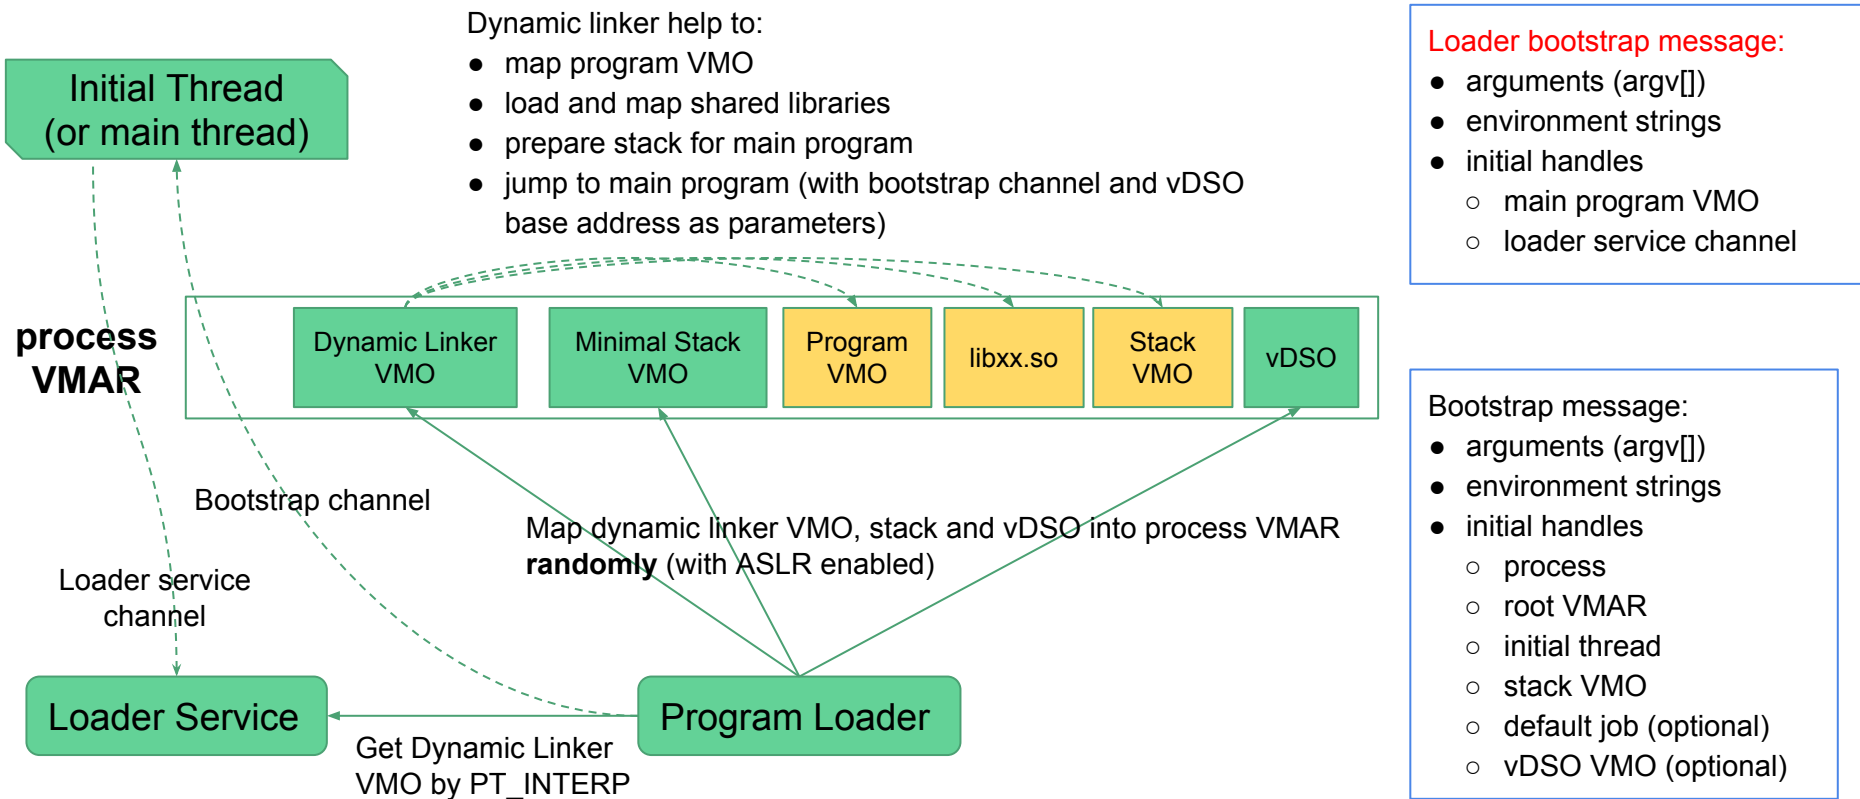
Program Loading and Dynamic Linking

- The kernel is not directly involved in normal program loading. (Except the userboot process which is the first userspace process booted by the kernel)
- Instead, the kernel merely provides the building blocks (VMO, process, VMAR, thread)
- Userspace environment use ELF format for machine-code executable files and provide a dynamic linker and C/C++ execution environment
- The main implementation of program loading resides in the **launchpad** library
- Program loading is based on VMOs and on IPC protocols used via channels
- Zircon only supports **P**osition-Independent **E**xecutables (ELF **ET_DYN** files)

Program Loading: ELF ET_DYN file with no PT_INTERP



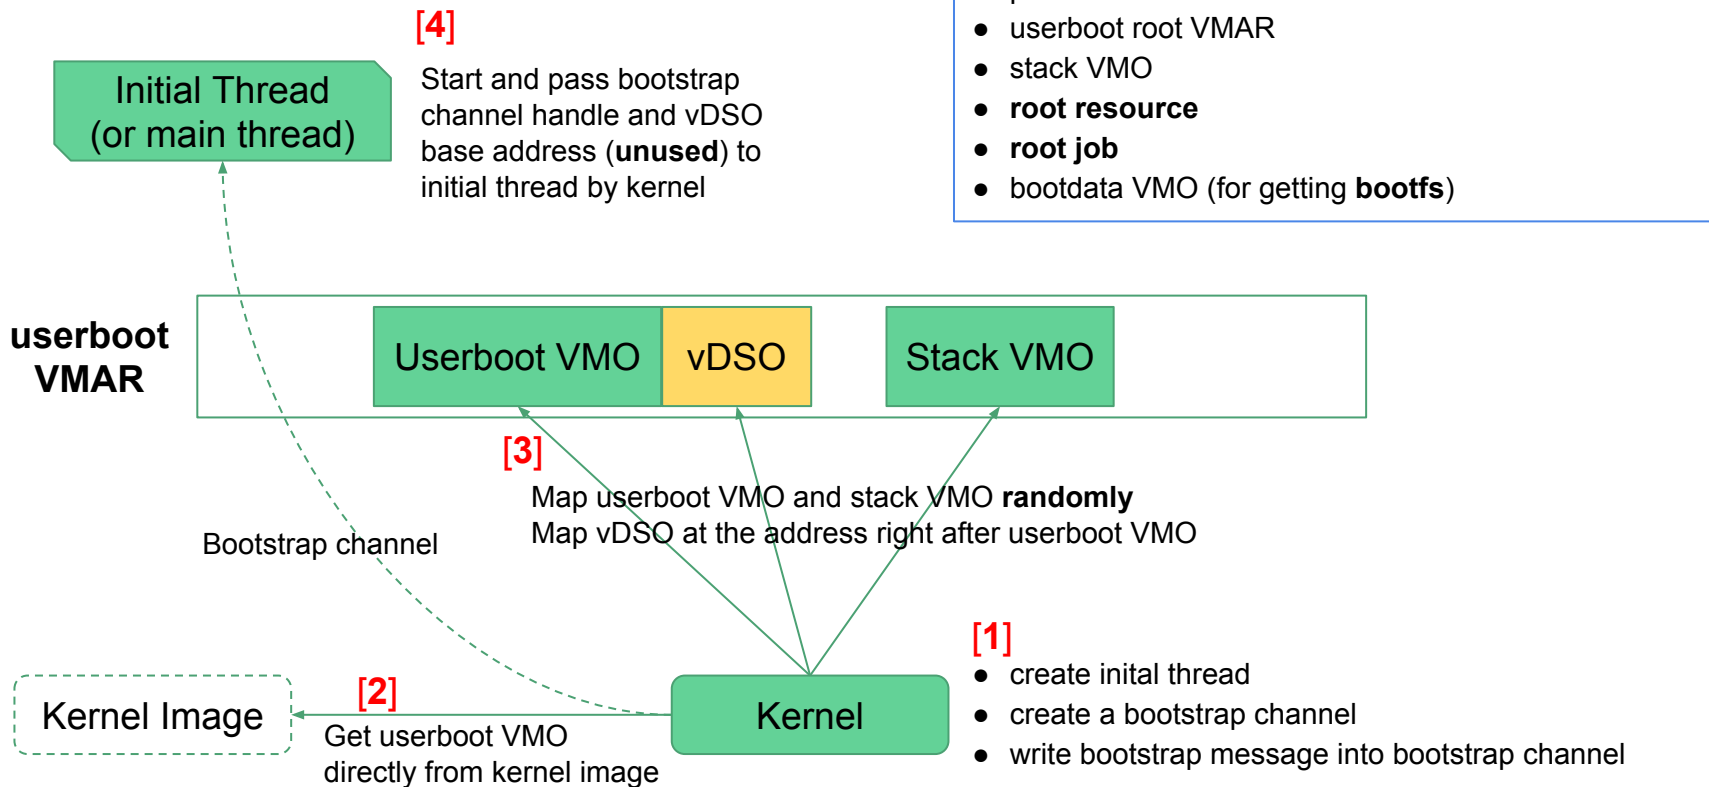
Program Loading: ELF ET_DYN file with PT_INTERP



Kernel to Userspace Bootstrapping

- userboot is the first userspace process started by the kernel
- userboot is built as an ELF dynamic shared object which embedded in the kernel image
- Due to the simple layout of userboot image, the kernel does not need to interpret its ELF header.
 - Information for loading userboot image will be extracted in compile time
 - Information including read-only segment size, executable segment size, entry point address
- In conclusion, the kernel starts the first user process **without** implementing:
 - ELF parser
 - Minimal file system
 - Decompression library

Userboot Startup



Scheduling

Scheduling

- Each CPU has its own set of priority queues levels from 0 to 31 (highest)
- Thread time slice rule:
 - If run out of its time slice, thread will be preempted and reinserted at the **end** of the appropriate priority queue
 - If still has time slice remaining, thread will be inserted at the **head** of the priority queue so it will be able to resume as quickly as possible
- Three different factors used to determine the effective priority of a thread:
 - **Base priority**: set at thread creation time or by `thread_set_priority()`
 - **Priority boost**: used to offset the base priority
 - **Inherited priority**: inherit from another higher priority thread which is blocked by the thread
- $\text{Effective_priority} = \text{MAX}(\text{base_priority} + \text{priority_boost}, \text{inherited_priority})$
- The intention is to ensure that **interactive threads are serviced quickly**

Realtime and Idle Threads

- Idle threads
 - Runs when no other threads are runnable
 - One on each CPU and lives outside of the priority queues
 - Used to track idle time for platform to implement the low power wait mode
- Realtime threads
 - Marked with `THREAD_FLAG_REAL_TIME`
 - Allowed to run **without preemption** and will run until they block, yield, or manually reschedule

Backup

Out-Of-Memory System (OOM)

- A kernel OOM thread is responsible for monitoring free memory size periodically
- OOM starts to work if the amount of free memory is lower than 'redline'
- OOM will pick and kill a job which is marked as 'kill_on_oom'
- OOM-ranker driver is under development

```
$ k oom info
[01725.664] 01100.01113> OOM info:
[01725.664] 01100.01113>   running: true
[01725.664] 01100.01113>   printing: false
[01725.664] 01100.01113>   simulating lowmem: false
[01725.664] 01100.01113>   sleep duration: 1000ms
[01725.664] 01100.01113>   redline: 50M (52428800 bytes)
```